

AWS Solutions Architect


Load Balancers, Auto Scaling Groups



Н. Ганжигүүр
Fibo Cloud

Proxy servers
Elastic Load Balancers
Auto Scaling Group
HA designs
Hands on

Proxy vs Reverse-proxy



1. Firewall
2. Better Management
3. Security
4. Caching
5. Encryption/Decryption

1. Protect servers
2. Caching
3. Compress

Stateless vs Stateful app

Stateless

Does not require the server to retain information about the state.

Server design, implementation and architecture is simple.

Handles crashes well, as we can fail over to a completely new server. Servers are regarded as cheap commodity machines.

Scaling architecture is easy.


Stateful

Requires a server to save information about a session.

Server design, implementation and architecture is complicated.


Does not handle crashes well. Servers are regarded as valuable and long-living. The user would probably be logged out and have to start from the beginning.

Scaling architectures is difficult and complex.



Elastic Load balancers

- ELB Stands for Elastic Load Balancer.
- It distributes the incoming traffic to multiple targets such as Instances, Containers, Lambda Functions, IP Addresses etc.
- It spans in single or multiple availability zones.
- It provides high availability, scaling and security for the application



- Classic Load Balancer (CLB) – this is the oldest of the three and provides basic load balancing at both layer 4 and layer 7.
- Application Load Balancer (ALB) – layer 7 load balancer that routes connections based on the content of the request.
- Network Load Balancer (NLB) – layer 4 load balancer that routes connections based on IP protocol data.

The Classic Load Balancer may be phased out over time and Amazon are promoting the ALB and NLB for most use cases within VPC.

Application Load Balancer

- It is best suited for load balancing of the web applications and websites.
- It routes traffic to targets within Amazon VPC based on the content of the request.

Network Load Balancer

- It is mostly for the application which has ultra-high performance.
- This load balancer also acts as a single point of contact for the clients.
- This Load Balancer distributes the incoming traffic to the multiple targets.
- The listener checks the connection request from the clients using the protocol and ports we specify.
- It supports TCP, UDP and TLS protocol.

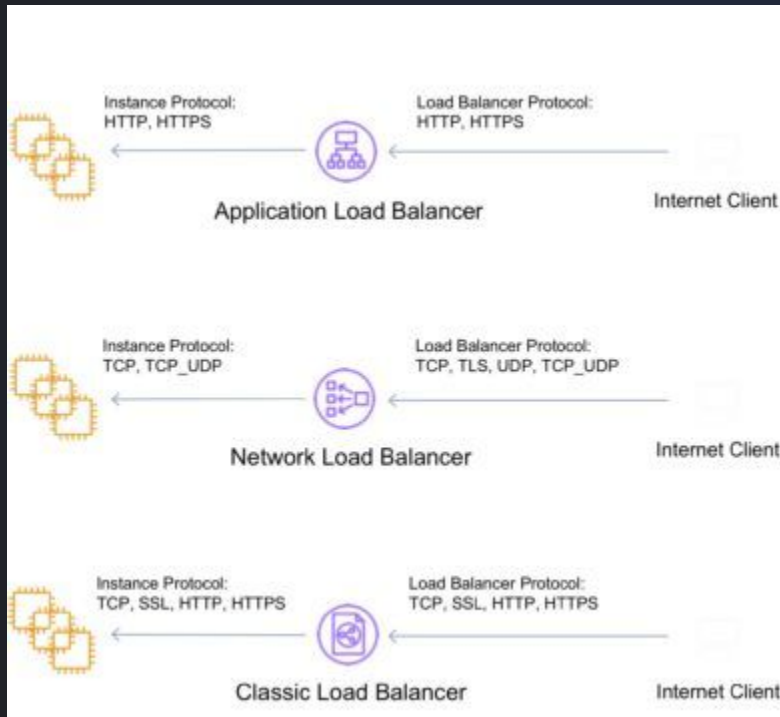
Gateway Load Balancer (Newly Introduced)

- It is like other load balancers but it is for third-party appliances.
- This provides load balancing and auto scaling for the fleet of third-party appliances.
- It is used for security, network analytics and similar use cases.

Classic Load Balancer

- It operates at request and connection level.
- It is for the EC2 Instance build in the old Classic Network.
- It is an old generation Load Balancer.
- AWS recommends to use Application or Network Load Balancer instead.

Types of ELB



Application Load Balancer

- Operates at the request level
- Routes based on the content of the request (layer 7)
- Supports path-based routing, host-based routing, query string parameter-based routing, and source IP address-based routing
- Supports IP addresses, Lambda Functions and containers as targets

Network Load Balancer

- Operates at the connection level
- Routes connections based on IP protocol data (layer 4)
- Offers ultra high performance, low latency and TLS offloading at scale
- Can have static IP / Elastic IP
- Supports UDP and static IP addresses as targets

Classic Load Balancer

- Old generation; not recommended for new applications
- Performs routing at Layer 4 and Layer 7
- Use for existing applications running in EC2-Classi

Feature	Application Load Balancer	Network Load Balancer	Classic Load Balancer
Protocols	HTTP, HTTPS	TCP	TCP, SSL, HTTP, HTTPS
Platforms	VPC	VPC	EC2–Classic, VPC
Health Checks	✓	✓	✓
CloudWatch Metrics	✓	✓	✓
Logging	✓	✓	✓
Zonal fail-over	✓	✓	✓
Connection draining	✓	✓	✓
Load balancing to multiple ports on an instance	✓	✓	
WebSockets	✓	✓	
IP addresses as targets	✓	✓	
Lambda functions as targets	✓		
Load balancer deletion protection	✓	✓	
Path-based routing	✓		
Host-based routing	✓		
HTTP header-based routing	✓		
HTTP method-based routing	✓		
Query string parameter-based routing	✓		
Source IP address CIDR-based routing	✓		
Native HTTP/2	✓		
Configurable idle connection timeout	✓		✓

Feature	Application Load Balancer	Network Load Balancer	Classic Load Balancer
Cross-zone load balancing	✓	✓	✓
SSL offloading	✓	✓	✓
Server Name Indication (SNI)	✓		
Sticky sessions	✓		✓
Back-end server encryption	✓	✓	✓
Static IP		✓	
Elastic IP address		✓	
Preserve source IP address		✓	
Resource-based IAM permissions	✓	✓	✓
Tag-based IAM permissions	✓	✓	
Slow start	✓		
User authentication	✓		
Redirects	✓		
Fixed response	✓		
Custom security policies			✓

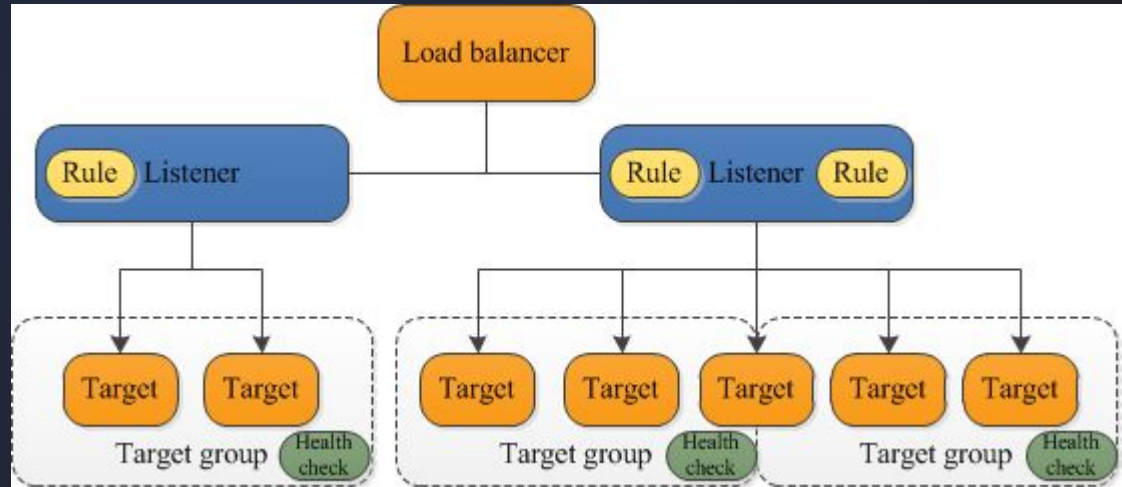
Listeners & Target Group

Listeners

- A listener is a process that checks for connection requests, using the protocol and port that you configured.
- You can add HTTP, HTTPS or both.

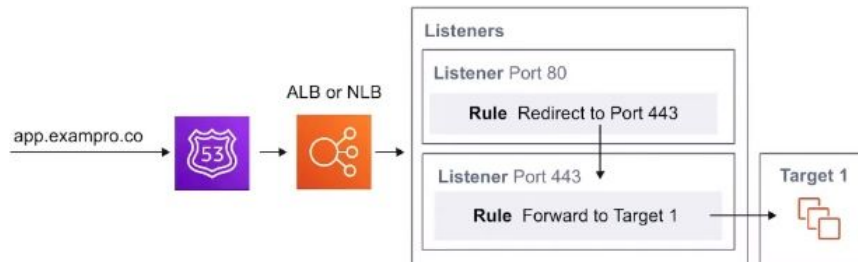
Target Group

- It is the destination of the ELB.
- Different target groups can be created for different types of requests.
- For example, one target group i.e., a fleet of instances will be handling the general request and other target groups will handle the other type of request such as micro services.
- Currently, three types of target supported by ELB: Instance, IP and Lambda Functions.



For Application Load Balancer (ALB) or Network Load Balancer (NLB) traffic is sent to the Listeners.

When the port matches it then checks the rules what do to. The rules will forward the traffic to a Target Group. The target group will evenly distribute the traffic to instances registered to that target group.



Description

Listeners

Monitoring

Integrated services

Tags

A listener checks for connection requests using its configured protocol and port, and the load balancer uses the listener rules to route requests to targets. You can add, remove, or update listeners and listener rules.


Add listener

Edit

Delete

Listener ID	Security policy	SSL Certificate	Rules
<input type="checkbox"/> HTTP : 80 arn...08e1a3165cec5d22 ▾	N/A	N/A	Default: redirecting to HTTPS://#{host}:443/#{path}?#{query} View/edit rules
<input type="checkbox"/> HTTPS : 443 arn...a64254b4731fa5c1 ▾	ELBSecurityPolicy	Default: a213d84c-4210-4fc7-3569-9113c39394fb (ACM) View/edit certificates	Default: forwarding to production View/edit rules

For Classic Load Balancer (CLB) traffic is sent to the Listeners. When the port matches it then it forwards the traffic to any EC2 instances that are registered to the Classic Load Balancer. CLB does not allow you to apply rules to listeners



Application Load Balancers are designed to balance **HTTP** and **HTTPS** traffic.

They **operate at Layer 7 (of the OSI Model)**.

ALB has a feature called **Request Routing** which allows you to add routing rules to your listeners based on the HTTP protocol.

Web Application Firewall (WAF) can be attached to ALB.

Great for Web Applications

OSI Layers

Layer 7 **Application**

Layer 6 **Presentation**

Layer 5 **Session**

Layer 4 **Transport**

Layer 3 **Network**

Layer 2 **Data Link**

Layer 1 **Physical**

Network Load Balancers are designed to balance **TCP/UDP**.

They **operate at Layer 4 (of the OSI Model)**

Can handle **millions of requests per second** while still maintaining extremely low latency.

Can preform Cross-Zone Load Balancing

Great for Multiplayer Video Games or When network performance is critical

OSI Layers

Layer 7 **Application**

Layer 6 **Presentation**

Layer 5 **Session**

Layer 4 **Transport**

Layer 3 **Network**

Layer 2 **Data Link**

Layer 1 **Physical**

It was AWS first load balancer (**legacy**)

Can balance **HTTP, HTTPS** or **TCP** traffic (not at the same time)

It can use **Layer 7-specific features (OSI Model)** such as **sticky sessions**.

It can also use **strict Layer 4 (OSI Model)** balancing for purely TCP applications.

Can preform Cross-Zone Load Balancing

It will respond with a **504 error (timeout)** if the underlying application is not responding. (**at the web-server or database level**)

Not recommended for use, instead use NLB or ALB

OSI Layers

Layer 7 **Application**

Layer 6 **Presentation**

Layer 5 **Session**

Layer 4 **Transport**

Layer 3 **Network**

Layer 2 **Data Link**

Layer 1 **Physical**

Web Application Deployed in Multiple Servers:

If a web Application/Website is deployed in multiple EC2 Instances then we can distribute the traffic between the Application Load Balancers.

Building a Hybrid Cloud:

Elastic Load Balancing offers the ability to load balance across AWS and on-premises resources, using a single load balancer. You can achieve this by registering all of your resources to the same target group and associating the target group with a load balancer.

Migrating to AWS:


ELB supports the load balancing capabilities critical for you to migrate to AWS. ELB is well positioned to load balance both traditional as well as cloud native applications with auto scaling capabilities that eliminate the guess work in capacity planning.

Charges:

- Charges will be based on each hour or partial hour that the ELB is running.
- Charges will also depend on the LCU (Load Balancer Units)
 - Number of new connections per second (up to 25 new connections per second is one LCU)
 - Number of active connections per minute (up to 3,000 active connections per minute is one LCU)
 - Bandwidth measured in Mbps (up to 2.22 Mbps is one LCU)

If you **need the IPv4 address** of a user, check the **X-Forwarded-For** header

The **X-Forwarded-For (XFF)** header is a command method for identifying the **originating IP address** of a client connecting to a web server through an HTTP proxy or a load balancer.



Sticky sessions

Sticky Sessions is an advanced load balancing method that allows you to **bind a user's session to a specific EC2 instance**.


Ensures all **requests** from that session are **sent to the same instance**.

Typically **utilized** with a **Classic Load Balancer**

Can be enabled for ALB though can only be set on a Target Group not individual EC2 instances.

Cookies are used to remember which EC2 instance.

Useful when specific **information is only stored locally on a single instance**




Cross-Zone Load Balancing

Only for **Classic** and **Network** Load Balancer

Cross-Zone Load Balancing **Enabled**

requests are distributed evenly across the instances **in all enabled** Availability Zones.

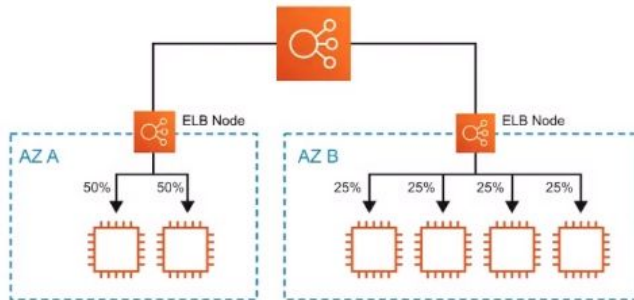


Attributes	
Deletion protection	Disabled
Cross-Zone Load Balancing	Disabled
Access logs	Disabled

[Edit attributes](#)

Cross-Zone Load Balancing **Disabled**

requests are distributed evenly across the instances **in only** its Availability Zone.



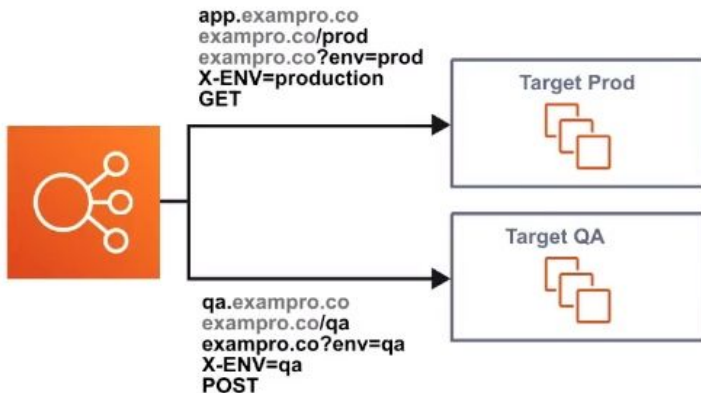
Edit load balancer attributes	
Delete Protection	<input type="checkbox"/> Enable
Cross-Zone Load Balancing	<input checked="" type="checkbox"/> Enable

Regional data transfer charges may apply when cross-zone load balancing is enabled. See the [documentation](#) for more information.

ALB forwarding

Apply rules to incoming request and then **forward** or **redirect** traffic.

- ✓ Host header
- ✓ Source IP
- ✓ Path
- ✓ Http header
- ✓ Http header method
- ✓ Query string



The screenshot shows the AWS IAM console interface for configuring rules. It displays two rule configurations. The top rule is for 'Target Prod' and the bottom rule is for 'Target QA'.

RULE ID	IF (all match)	THEN
1	<div>+ Add condition</div> <div>Host header...</div> <div>Path...</div> <div>Http header...</div> <div>Http request method...</div> <div>Query string...</div> <div>Source IP...</div>	<div>+ Add action</div> <div>Forward to production</div>

last HTTPS 443: default action
This rule cannot be moved or deleted

RULE ID	IF (all match)	THEN
1	<div>+ Add condition</div>	<div>+ Add action</div> <div>Forward to...</div> <div>Redirect to...</div> <div>Return fixed response...</div> <div>Authenticate...</div> <div>Forward to production</div>

last HTTPS 443: default action
This rule cannot be moved or deleted

IF	THEN
✓ Requests otherwise not routed	Forward to production


Auto Scaling Group

AWS Auto-Scaling Group

Auto Scaling Groups (ASG) contains a collection of EC2 instances that are treated as a group for the purposes of automatic scaling and management

Concepts:

- Group (fleet)
 - Logical component. Webserver group or Application group or Database group etc.
- Configuration template
 - Groups uses a launch template or launch config as a config template for its EC2 instance. AMI ID, instance type, keypair, security groups, block device mappings...
- Scaling Policy
 - Scaling options provides several ways for you to scale your Auto Scaling groups.
- Health check



AWS ASGs


Scaling policies

1. Maintain current instance levels at all times

You can configure your ASG to maintain a specified number of running instance at all times.
When Amazon EC2 Auto scaling finds an unhealthy instance, it terminates it and launches a new one.

2. Scale manually

Manually scaling is the most basic way to scale your resources, where you specify only the change in the min, max or desired capacity of your ASG.



3. Scale based on a schedule

Scaling by schedule means that scaling actions are performed automatically as a function of time and date.

This is useful when you know exactly when to increase or decrease the number of instances in your group, simply because the need arises on a predictable schedule.

4. Scale based on demand

A more advanced way to scale your resources. Using scaling policies lets you define parameters that control that scaling process. Ex: CPU load > 75%

5. Use predictive scaling

Combination with ASG to scale resource across multiple services.

Predictive scaling forecasts and scheduled actions [Info](#)

Predictive scaling mode

Forecast and scale

Load metric

Total CPU utilization

Forecast period

2 days

Max capacity behavior

Set forecast capacity to max capacity


Forecast granularity

60 minutes


Scheduled action buffer time

300 seconds

Load (Total CPU utilization)



Capacity (Instance count)



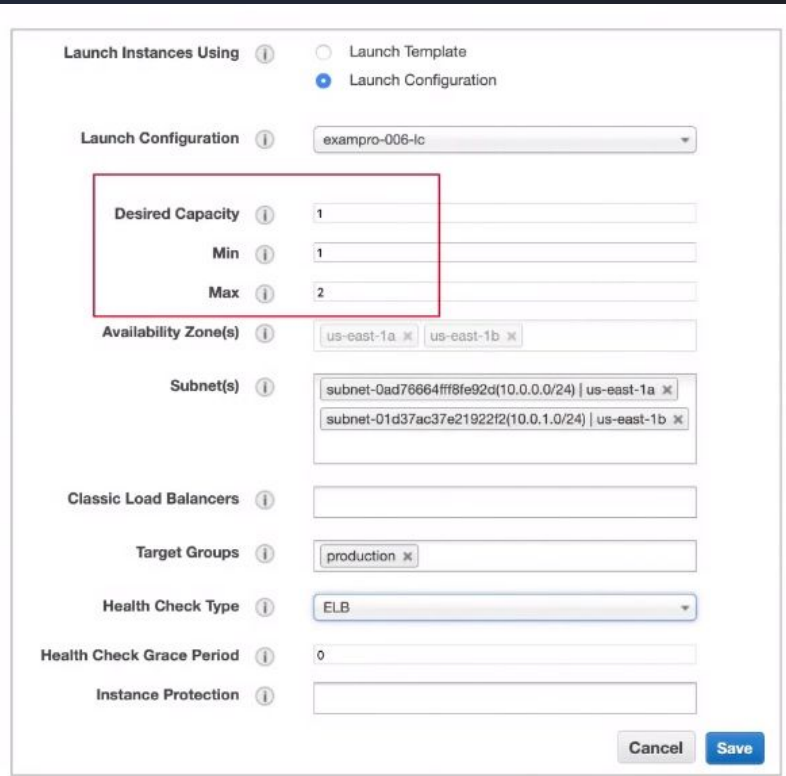
The size of an Auto Scaling Group is based on **Min**, **Max** and **Desired Capacity**.


Min is how many EC2 instances should at least be running.


Max is number EC2 instances allowed to be running.


Desired Capacity is how many EC2 instances you want to ideally run.


ASG will always launch instances to meet minimum capacity.





Launch Instances Using  ☐ Launch Template ☒ Launch Configuration


Launch Configuration  exampro-006-lc


Desired Capacity  1


Min  1


Max  2


Availability Zone(s)  us-east-1a x us-east-1b x


Subnet(s)  subnet-0ad76664fff8fe92d(10.0.0.0/24) | us-east-1a x
subnet-01d37ac37e21922f2(10.0.1.0/24) | us-east-1b x

Classic Load Balancers 

Target Groups  production x

Health Check Type  ELB

Health Check Grace Period  0


Instance Protection 

Cancel Save

EC2 Health Check Type

ASG will perform a health check on EC2 instances to determine if there is a software or hardware issue. This is based on the **EC2 Status Checks**. If an instance is considered unhealthy, ASG will terminate and launch a new instance.


✓ 2/2 checks passed



Health Check Type ⓘ	EC2
Health Check Grace Period ⓘ	EC2
Instance Protection ⓘ	ELB

ELB Health Check Type

ASG will perform a health check based on the ELB health check. ELB can perform health checks by pinging an HTTP(S) endpoint with an expected response. If ELB determines a instance is unhealthy it forwards this information to ASG which will terminate the unhealthy instance.



Health Check Type ⓘ	EC2
Check Grace Period ⓘ	EC2
	ELB

Scaling Out: Adding More Instances

Scaling In: Removing Instances

Target Tracking Scaling Policy

Maintains a specific metric at a target value.

eg. If **Average CPU Utilization** exceeds 75% then add another server.

Create Scaling policy

Name:

Metric type: Application Load Balancer Request Count Per Target
✓ Average CPU Utilization
Average Network In (Bytes)
Average Network Out (Bytes)

Target value:

Instances need: seconds to warm up after scaling

Disable scale-in: ☐

Simple Scaling Policy

Scales when an **alarm is breached**.

Create Scaling policy

Name:

Execute policy when:

No alarm selected

Take the action:

Add

0

instances

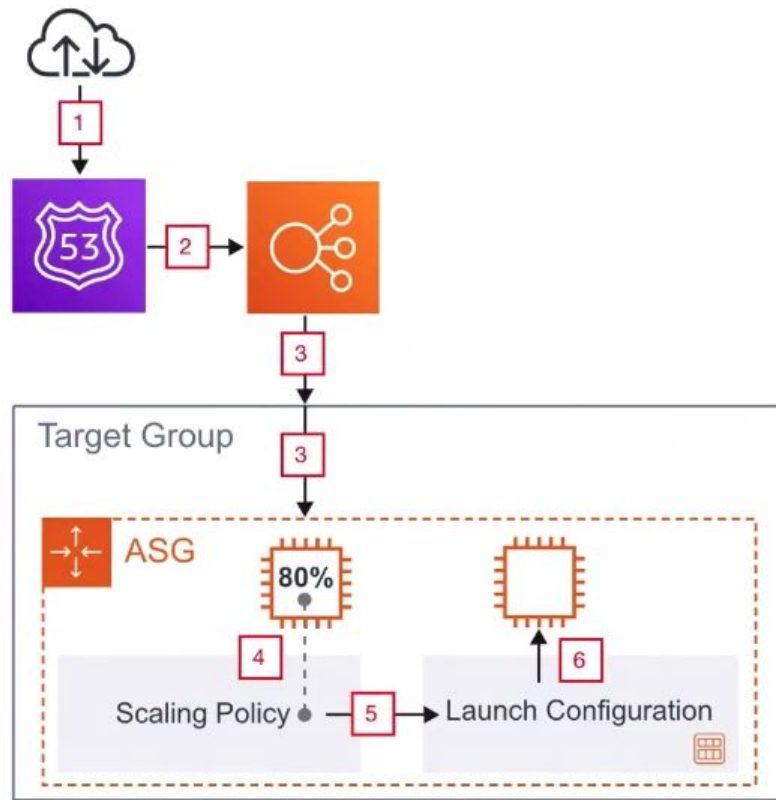
And then wait:

300

seconds before allowing another scaling activity

Not recommended, legacy scaling policy. Use scaling policies with steps now.

1. Burst of traffic from the internet hits our domain.
2. Route53 points that traffic to our load balancer.
3. Our load balancer passes the traffic to its target group.
4. The target group is associated with our ASG and sends the traffic to instances registered with our ASG
5. The ASG Scaling Policy will check if our instances are near capacity.
6. The Scaling Policy determines we need another instance, and it Launches an new EC2 instance with the associated Launch Configuration to our ASG



A launch configuration is an instance configuration template that an Auto Scaling group uses to launch EC2 instances.

AUTO SCALING
Launch Configurations
Auto Scaling Groups

Launch Configuration ⓘ example-007

A Launch Configuration is the same process as Launching an EC2 instance except you are saving that configuration to Launch an Instance for later. Hence “Launch Configuration.

1. Choose AMI 2. Choose Instance Type 3. Configure details 4. Add Storage 5. Configure Security Group 6. Review

Create Launch Configuration

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch an EC2 instance; or you can select one of your own AMIs.

Quick Start

- My AMIs
- AWS Marketplace
- Community AMIs



Amazon Linux
Free tier eligible

Amazon Linux 2 AMI (HVM), SSD Volume Type - ami-0b898040803850f

Amazon Linux 2 comes with five years support. It provides Linux kernel 4.14 tuned for 2.26, Binutils 2.29.1, and the latest software packages through extras.

Root device type: ebs Virtualization type: hvm

Launch Configurations **cannot be edited**, When you need to update your Launch Configuration you create a new one or clone the existing configuration and then manually associate that new Launch Configuration

Launch Templates are Launch Configurations with Versioning, Everyone appears to still use Launch Configurations

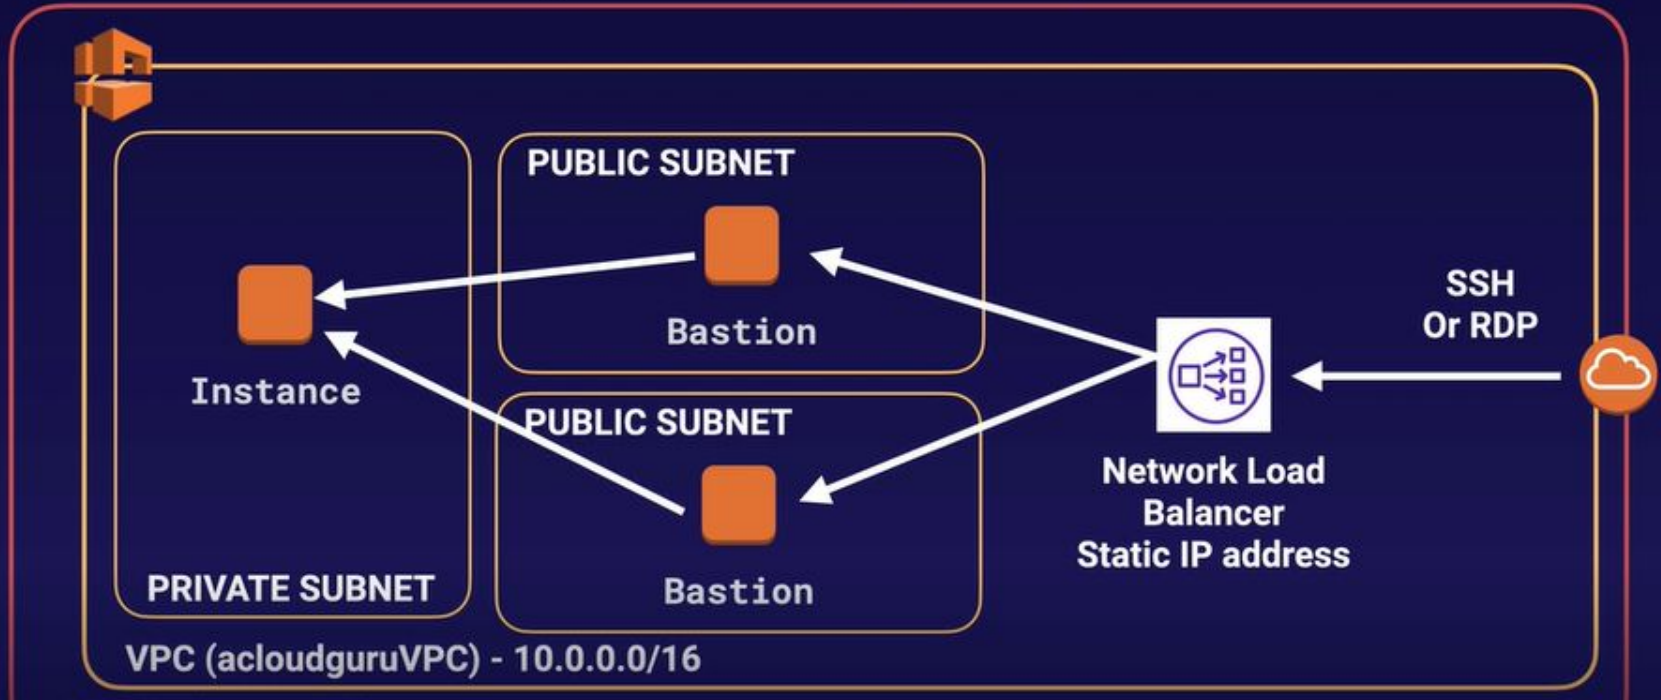
```
#!/bin/bash
yum update -y
yum install httpd -y
systemctl start httpd
systemctl enable httpd
EC2ID=$(curl -s http://169.254.169.254/latest/meta-data/instance-id)
EC2AZ=$(curl -s http://169.254.169.254/latest/meta-data/placement/availability-zone)
echo '<html><h1>My EC2: INID  AZ: AZID</h1></html>' > /var/www/html/index.txt
sed -e "s/AZID/$EC2AZ/g" -e "s/INID/$EC2ID/g" /var/www/html/index.txt >
/var/www/html/index.html
```

```
sudo amazon-linux-extras install epel -y
sudo yum install stress -y
```

HA Bastions Hosts

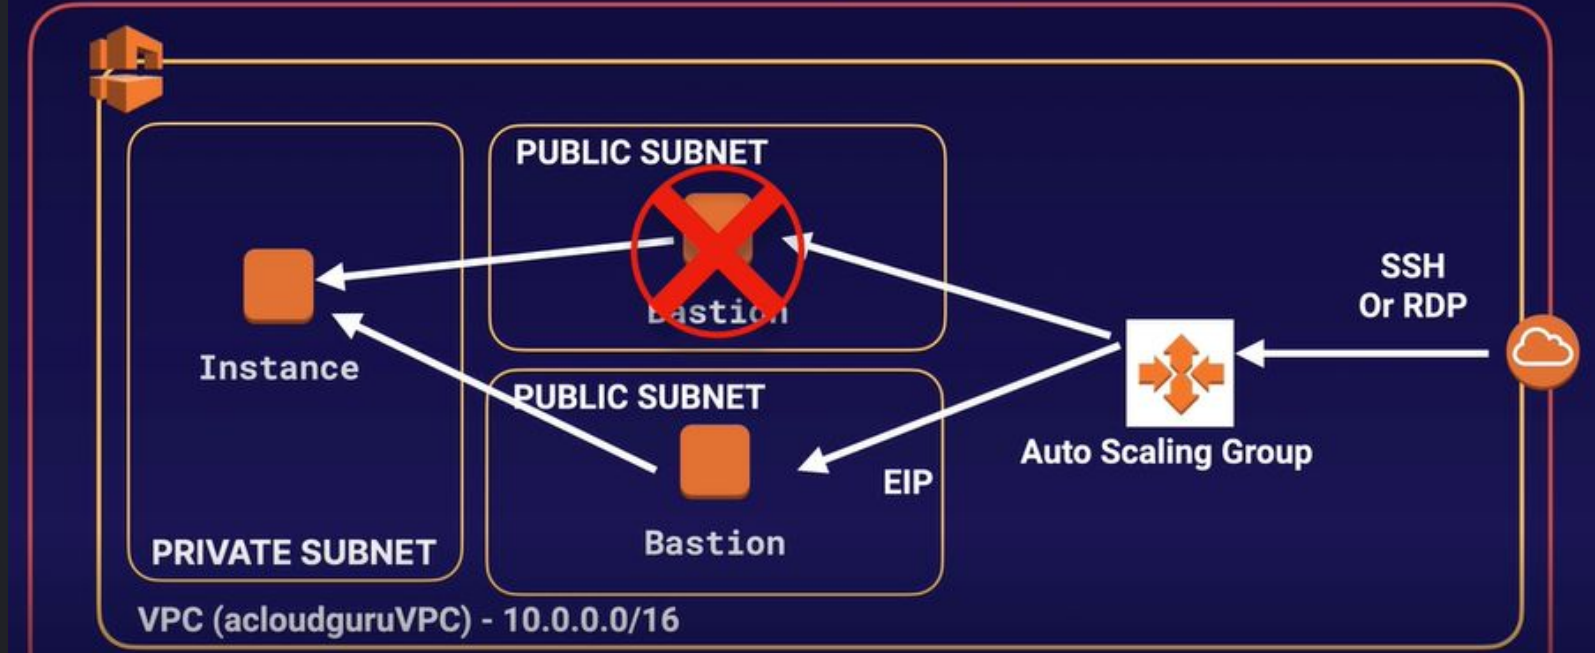
Scenario 1:

Two EC2 Instances, Two Availability Zones, Network Load Balancer



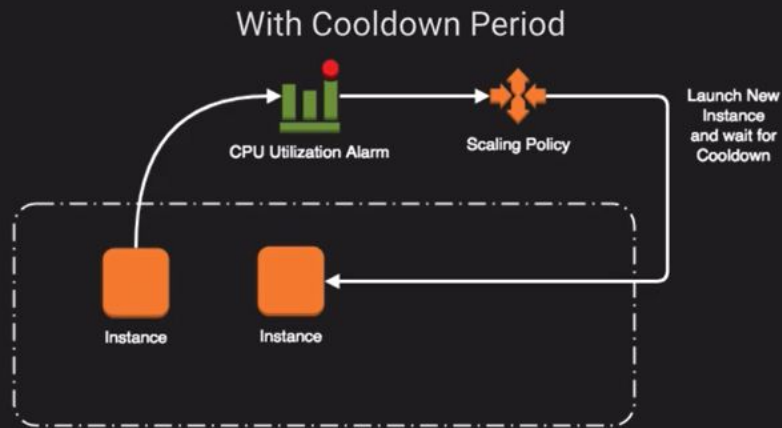

HA Bastions Hosts

Scenario 2:
One EC2 Instance, Two Availability Zones, Auto Scaling Group



Cooldown Periods

- Configurable duration that gives your scaling a chance to “come up to speed” and absorb load.
- Default cooldown period is 300 seconds.
- Automatically applies to **dynamic scaling** and optionally to manual scaling but **not supported for scheduled scaling**.
- Can override default cooldown via scaling-specific cool down



Scaling with SQS

